

EPIC-I — Research/News Feed — Technical Spec

Status: Draft — for stakeholder review **Date:** 14 July 2026 **Author:** Adrian Hall (Technical Lead), drafted with Claude Code **Scope:** The ninth and final per-epic technical spec. Unlike EPIC-A through EPIC-H, this epic already has substantial design work done — the architecture spec's Section 8 covers source adapters, ingestion, deduplication, classification, and scoring in real detail, validated by the working prototype in `prototypes/research-feed/`. This spec is deliberately lighter: it consolidates what's already decided, adds the member-facing API surface the architecture spec didn't specify, and carries forward the open items already identified there rather than re-deriving them.

Source of truth: the architecture spec's Section 8 (`docs/superpowers/specs/2026-07-14-askapeer-architecture-design.md`) and `docs/AHP_Research_Feed_Design_Conversation.md/docs/research-feed-sources-and-roadmap.md`. **Reminder:** EPIC-I is not yet in the PRD's own Section 6.1 MoSCoW list — it was added during architecture design and still needs reflecting back to Paul Gouge and Andrew Renshaw (architecture spec, Section 11 and Section 1).

Contents

1. Scope
 2. Data model (as already specified)
 3. Member-facing interest management
 4. member_interests vs. the missing tag-follow feature
 5. Taxonomy dependency on EPIC-C
 6. API endpoints
 7. Carried-forward open items
 8. Test plan
 9. Open questions
-

1. Scope

In scope: source adapters and ingestion (already designed), deduplication and classification (already designed and prototype-validated), the feed API, and — the actual new content in this spec — how a member manages their interest profile, and the relationship between that profile and forum tag-following.

Out of scope: everything the architecture spec's Section 8 already fully specifies is referenced, not repeated, here.

2. Data model (as already specified)

Reproduced from the architecture spec, Section 8, for reference only — no changes proposed:

```
research.articles          (id, doi, pmid, other_ids, title,
abstract, journal,
                           published_date, evidence_type,
open_access, source_refs, created_at)
research.ingestion_cursors (source_name, last_run_at, last_cursor)
community.member_interests (handle_id, tag, weight, updated_at)
```

Source adapters (fetchSince(cursor): RawArticle[]), deduplication order (DOI → PMID → other IDs → normalised title/year → fuzzy match), rule-based classification, and precomputed scoring (topic match + evidence-type weighting + recency) are all as specified there.

3. Member-facing interest management

The architecture spec designs the ingestion and scoring pipeline in detail but doesn't specify how a member actually builds their `member_interests` profile — the working prototype does this via “interest tags on the Profile pane” (per docs/research-feed-sources-and-roadmap.md), which this spec formalises as an API:

```
GET  /v1/research-feed/interests
PUT  /v1/research-feed/interests { tag, weight } -- upsert;
weight e.g. 0-1 or a small
```

integer scale, matching whatever

the

prototype's scoring already

expects — not redefined here

```
DELETE /v1/research-feed/interests/:tag
```

`member_interests` is handle-scoped (architecture spec, Section 8's data-boundary note) — this is ordinary `community-schema` data, no identity involvement, consistent with everything else in that schema.

4. `member_interests` vs. the missing tag-follow feature

Worth flagging precisely because it's easy to conflate:

`community.member_interests` (this epic, weighted, used to score research-literature relevance) and the tag-follow mechanism EPIC-G's spec identified as missing (its Section 7, needed for the forum's Should-have “personalised feed” and “email digest”) are **two different concepts that happen to both be “a member's interest in a tag.”** One is a weighted relevance signal for an external content stream (research articles); the other would be a binary follow relationship for internal content (forum posts). They could plausibly be

unified into a single `community.tag_interests`-style concept feeding both the research feed and the forum's personalised feed — or they could legitimately stay separate given they serve different content types with different scoring needs. This is a real design choice worth making deliberately once EPIC-C/EPIC-G's gap is addressed, not something this spec should resolve unilaterally by either building a duplicate mechanism or assuming unification. Flagged in Section 9 and cross-referenced from EPIC-G's spec.

5. Taxonomy dependency on EPIC-C

EPIC-C's spec (Section 3) already flags that the forum's category/tag taxonomy, Andrew Renshaw's body-area list, and the research feed prototype's `taxonomy.json` are three overlapping-but-not-unified vocabularies. This epic depends on that reconciliation directly: `member_interests.tag` and the classification taxonomy the architecture spec's Section 8 describes should end up being **the same controlled vocabulary** as whatever EPIC-C settles on for forum tags, not a separately-maintained list — otherwise a member's forum tag-following and their research-feed interests would use different, possibly inconsistent vocabularies for what a member experiences as “topics I care about.” No new work is proposed here beyond restating this dependency clearly, since EPIC-C's spec already owns resolving it.

6. API endpoints

```
GET /v1/research-feed?cursor=...      -- already specified,
architecture spec Section 8:
                                     combines
member_interests weights with precomputed
                                     article scores, returns
an explanation string

GET /v1/research-feed/interests      -- Section 3, new in this
spec
PUT /v1/research-feed/interests
DELETE /v1/research-feed/interests/:tag
```

7. Carried-forward open items

These are already identified in the architecture spec's Section 11 as EPIC-I pre-launch items — restated here as this epic's responsibility, not re-analysed:

- **DOAJ integration** (predatory-journal legitimacy check) and **Retraction Watch / Crossref integration** (flagging withdrawn papers) — sources identified, integration not yet built.
- **PEDro API access** — worth a direct enquiry to the PEDro team (University of Sydney/NeuRA) given it's the best domain-specific fit found so far, but no confirmed public API.

- **Abstract-redistribution licensing review** — unresolved, relevant if a commercial source (Elsevier/Scopus/Cochrane) is ever added.
-

8. Test plan

- **Interest CRUD:** creating/updating/deleting a `member_interests` row is handle-scoped and has no `identity` involvement (same access-boundary test pattern as every other epic's spec).
 - Everything else (deduplication, classification, scoring) is already covered by the working prototype's own validation, per the architecture spec — this spec doesn't add new test surface beyond the member-facing interest endpoints.
-

9. Open questions

- **PRD update:** EPIC-I still needs adding to Section 6.1's MoSCoW list — a standing item from the architecture spec, restated here since it's this epic specifically that's affected.
- `member_interests` **vs. tag-follow unification** (Section 4): a genuine design choice spanning this epic, EPIC-C, and EPIC-G — needs a decision once the tag-follow gap (EPIC-G's Section 7) is addressed.
- **Taxonomy reconciliation** (Section 5): depends entirely on EPIC-C's Section 3 resolution — this epic has no independent path to resolving it.
- **The three carried-forward items** (Section 7): DOAJ/Retraction Watch integration, PEDro enquiry, and licensing review all remain open pre-launch work, unchanged from the architecture spec's own assessment.